

RevenueOS


One seam, closed end to end.

Sales wins a deal. PreSales has two business days to start the POC plan.

In the US it takes 6.8. This is the loop that closes the gap.

The seam

Marketing → Sales → PreSales → Delivery → Support. It stalls at one join.



Not a headcount problem. The templates already exist — nothing fetches them.

The loop

Five steps. One screen. The human presses the last button.



Automated up to the commitment. The commitment stays with a named human.

One source of truth

Nothing on the screen is stored. Every number is derived from an append-only log.

Event log

brief.arrived
draft.generated
draft.accepted
draft.rejected
sla.breached
trigger.fired

age → business days since arrival

status → green / amber / red

reuse rate → accepted ÷ decided

boosts → ±8 per decision

triggers → who was escalated

Every number on screen has an ⓘ that shows the query behind it. If it cannot, it does not go on the screen.

Retrieval you can argue with

A score a Solution Architect can push back on beats a similarity number they cannot.

Same segment	30
Same regulator / analogue	16 / 8
Problem overlap	≤ 26
Systems overlap	≤ 18
Recently used	5
Proven reuse	4
Learned from decisions	± 8

Threshold 40

Nothing below it is shown.


No match = the plan says so.

Each match ships the reasons it fits, in words.

At ~200 templates I put embeddings in front for recall and keep this for the explanation.

The agent

Three nodes, one job. It drafts; it does not decide.



If the model is slow or unavailable, it falls back to a deterministic plan — and the screen says “sample draft” instead of “drafted by AI model”.

The demo never dies, and it never pretends.

Ask the seam

A chat box with no memory and six read-only tools.

get_events

get_health

get_brief

get_template

explain_match

read_doc

“Which briefs are past SLA right now and who was triggered?” → answers from the log, and names the sources it read.

“What is the weather in Mumbai?” → refuses in one sentence. No tool holds it, so there is no answer.

It learns from the decision, not from a rating

No feedback form. The button is the signal.

Accept → + 8 to every template that draft used

Reject → - 8 to every template that draft used

Boosts are derived from the event log — the server reads them itself and ignores anything the browser sends.

Refresh the page and the decision is still there. It was never browser state.

What is real, and what breaks first

Said before anyone has to ask.

Real

Event log, derive functions, tests

Retrieval scoring and its reasons

The agent, the tools, the refusals

SLA breach detection and escalation

Not real yet

Data is synthetic and labelled on screen

Store is in memory — a cold start loses it

No auth, one tenant, no workers

Breaks first at 10×: the store (#53)

48 open issues across six milestones. Nothing is closed unless code exists behind it.

The one thing I would never hand to the agent

The accept.

Pressing Accept is not a document being finished. It is a promise to a customer about what we will prove, in how many weeks, under their regulator.

The person who makes that promise has to be the person who will be in the room when it slips.